

When “Justice” Is Written in Code

AI, Epistemic Injustice, and Palantir’s ImmigrationOS

Claire Vlases

In April, Immigration and Customs Enforcement (ICE) paid software company Palantir thirty million dollars to build an “Immigration Lifecycle Operating System,” known as ImmigrationOS, an AI system that helps agents identify people for removal from the United States. ICE renewed the contract in September, bringing the total up to around sixty million dollars [5,6]. ImmigrationOS is one of many AI tools becoming rapidly integrated into the public sector. The Department of Homeland Security reports around 200 different uses of AI across its agencies [1].

ICE Is Paying Palantir \$30 Million to Build ‘ImmigrationOS’ Surveillance Platform

However, emerging research suggests that the use of AI in border and migration control has noticeably harmful effects. A recent study from the British Institute of International and Comparative Law (BIICL) found that AI in immigration builds discriminatory structures that continue to operate long after the initial deployment [13].

ICE adopted ImmigrationOS as a part of a campaign to identify and detain the people it calls the “Worst of the Worst” [6]. But, it is unclear how it exactly classifies people as the “worst.” ICE and Palantir have declined to disclose how many people the system tracks, which agencies it pulls data from, and whether there are safeguards to protect people from mistaken identity or excessive surveillance [6].

The warning is clear:


“Shrouded in secrecy and complexity, algorithmic decisions might well perpetuate bias and prejudice” [4].

From what we do know, the ImmigrationOS tool lets ICE search people using hundreds of micro-categories, like how they entered the US, their hair and eye color, scars, tattoos, and even detailed license-plate reader data showing everywhere someone drives.... all packaged as ‘objective’ data [10].

AI systems present themselves as objective, but they encode human judgments, biases, and power. Who gets to define what justice looks like when it’s embedded in code?

Palantir insists, “we build the tools, not the rules” [8]. As AI becomes increasingly embedded into our systems, I caution against this language. Choices about which variables matter, which features to track, and which patterns to prioritize are normative decisions.

We have seen the danger of algorithmic-overreliance before. More than a decade ago, ProPublica’s investigation into the Correctional Offender Management Profiling for Alternative Sanctions (COMPAS) tool, used to assess recidivism risk, revealed noticeable racial disparities [3].



The research found that the formula was particularly likely to falsely flag black defendants as future criminals, wrongly labeling them as high-risk at almost twice the rate as white defendants [3].

These were not technical errors. Rather, they were design decisions about which factors to weigh, what ‘risk’ means, and who gets to define it. The COMPAS recidivism algorithm is a poignant example of how automated decision systems that *appear* neutral can have harmful consequences.

ImmigrationOS belongs to the same family of systems. It reinforces a pattern in which algorithmic tools are introduced into high stakes decisions without adequate scrutiny of their epistemic foundations.

AI algorithms can contribute to such injustices because they are epistemic technologies; they consume, curate and produce content that guides greater understanding. It’s easy to simplify the problem, arguing that AI is only as good as the data on which it is trained, but there is much more to it. These AI systems don’t just classify, they generate knowledge-like content. AI intervenes directly in how knowledge is produced. It participates in inquiry, organizes and transforms data, and produces outputs through inferential and predictive operations that mimic forms of reasoning, and institutions treat AI generations as meaningful.

When programs make decisions, they exert power through their contributions to knowledge [9]. By shaping understanding, these algorithms have the capacity to do epistemic harm.

“Algorithmic decision procedures can reproduce existing patterns of discrimination, inherit the prejudice of prior decision makers, or simply reflect the widespread biases that persist in society... It can even have the perverse result of exacerbating existing inequalities by suggesting that historically disadvantaged groups actually deserve less favorable treatment” [12].


A deeper philosophical problem emerges. The opacity of AI systems facilitates a form of *epistemic injustice*: people are wronged in their capacity as knowers. The ignorance about a technology’s inner workings complicates the possibility of contesting its unjust decisions, and the relevant interpretive framework is controlled by an inscrutable technical structure rather than by shared social understanding [11]. When the logic of a system remains hidden, individuals cannot access the reasons used to classify them, nor can they supply counter reasons of their own. The system appears to know, but it refuses to explain.

This is a case of *hermeneutical injustice*; a type of epistemic injustice where individuals or groups lack the shared concepts or tools to understand and articulate their own social experiences, often because those experiences are marginalized by dominant social norms [9].

To understand the concept of hermeneutical injustice, consider a classic example: Survivors of sexual assault, before the term had real public recognition (in the 1970s), lacked a widely accepted concept to understand and articulate what had happened to them. The experience remained, but the conceptual tools were missing from the dominant repertoire.

AI systems can help create similar gaps, but in a more subtle way. As Google Deepmind researcher Jackie Kay writes, “the repository of collective understanding bears the imprints of dominant groups, leaving the experiences of marginalized groups misrepresented or distorted” [9]. When large scale models generate new classifications without input from marginalized communities, they update it in a direction that reinforces those same power relations.

Kay proposes the term '*generative hermeneutical ignorance*' to describe how AI systems can erase or misportray marginalized groups. Generative models lack the nuanced comprehension of human experience necessary to produce accurate and equitable representation. When state institutions adopt these outputs, they stabilize those distortions as the *official way* of knowing. We then act on these outputs as though the system had discovered a pattern that was already there, rather than one that it helped to construct [9]. Relying on the very instruments that uphold these structures perpetuates inequity, an ouroboros of injustice.



This produces a layered injustice. First, a system can mischaracterize people because it lacks the concepts and context required to fully capture the human experience. Second, the people who are mischaracterized have no real avenue to correct the record. They cannot access the model, examine its assumptions, or propose alternative categories. The system has authority, but not accountability.

Models like ImmigrationOS make these abstract concerns concrete. Within such a system, a person becomes a cluster of data points. The model arranges these features into a narrative. Because the system shields its logic behind proprietary protections and technical complexity, individuals cannot ask which variables mattered, why certain patterns were interpreted as suspicious, or how the model would respond if their circumstances changed. They are denied the very tools that would allow them to articulate objections in terms the system would recognize. This is hermeneutical injustice in a strict sense. It is a denial of the capacity to participate in the shared construction of meaning about one's own life.

Despite these concerns, this is not to say that we should do away with AI or refuse the use of algorithms in human decision-making. There is research that shows the opposite. Algorithmic systems can reduce human bias by standardizing decision procedures and exposing discriminatory patterns that would otherwise remain invisible. Significant efforts now exist to mitigate bias in algorithms, including an explosion of work on fairness metrics [2]. Cass Sunstein, a legal scholar, even argues that algorithms can be designed to avoid racial or other unlawful forms of discrimination. [16]. These points suggest that algorithmic tools have the potential to improve upon human judgment rather than undermine it.

Importantly, none of these counterarguments undermine the claim that AI systems can create an algorithmic hermeneutical injustice. Scholars who defend algorithmic systems do not argue that standardization itself produces fairness; it can reduce bias only when the underlying categories and interpretive resources are already just, inclusive, and publicly contestable. Sunstein's suggestion that algorithms can avoid unlawful discrimination presupposes that the categories the system uses have been constructed through transparent and participatory processes. Fairness metrics also require scrutinizing the model's assumptions.

Noam Chomsky's reminder that AI does not *understand* meaning is important here. These systems manipulate symbols according to statistical rules. They do not possess reasons, intentions, or moral concepts; rather, they "trade merely in probabilities that change over time" [5]. For that reason, any appearance of understanding comes from the human and institutional context that interprets the outputs as meaningful. The moral responsibility rests entirely with those who choose to grant these systems interpretive authority over human lives.



“The human mind is not, like ChatGPT and its ilk, a lumbering statistical engine for pattern matching, gorging on hundreds of terabytes of data and extrapolating the most likely conversational response...” [5]

When AI generates new interpretive structures, the responsibility lies with the humans and institutions who design the systems. The tools inherit the biases and blind spots of the institutions that deploy them.

As visionary poet Audre Lorde warns, “the master’s tools will never dismantle the master’s house” [14]. We cannot expect justice to emerge from algorithmic systems built on unjust foundations. Systems that reshape the conceptual resources of entire communities, while concealing their own standards of reasoning erode the conditions under which people can understand and challenge power and threaten the conditions that make democratic accountability possible.



Works Cited

- [1] "AI Use Case Inventory Library." U.S. Department of Homeland Security, 1 July 2025, www.dhs.gov/publication/ai-use-case-inventory-library.
- [2] Alex Beutel, Jilin Chen, Tulsee Doshi, Hai Qian, Allison Woodruff, Christine Luu, Pierre Kreitmann, Jonathan Bischof, and Ed H. Chi. 2019. Putting Fairness Principles into Practice: Challenges, Metrics, and Improvements. In Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society (AIES '19). Association for Computing Machinery, New York, NY, USA, 453–459. <https://doi.org/10.1145/3306618.3314234>
- [3] Angwin, Julia, et al. "Machine Bias." ProPublica, ProPublica, 23 May 2016, www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing.
- [4]: Cerquitelli, T., Quercia, D., & Pasquale, F. (Eds.). (2017). Transparent Data Mining for Big and Small Data (Vol. 32). Springer International Publishing. Preface, p. v.
- [5] Chomsky, Noam, et al. "Opinion: Noam Chomsky: The False Promise of Chatgpt." Noam Chomsky: The False Promise of ChatGPT, The New York Times, 8 Mar. 2023, web.archive.org/web/20230429120246/https://www.nytimes.com/2023/03/08/opinion/noam-chomsky-chatgpt-ai.html.
- [6] Dou, E. (2025, December 3). How Palantir shifted course to play key role in ICE deportations. The Washington Post. <https://www.washingtonpost.com/technology/2025/12/03/palantir-immigration-ice/>
- [7] Haskins, Caroline. "Ice Is Paying Palantir \$30 Million to Build 'ImmigrationOS' Surveillance Platform." Wired, Conde Nast, 18 Apr. 2025, www.wired.com/story/ice-palantir-immigrationos/.
- [8] Hubbard, Steven. "Ice to Use ImmigrationOS by Palantir, a New AI System, to Track Immigrants' Movements." American Immigration Council, American Immigration Council, 21 Aug. 2025, www.americanimmigrationcouncil.org/blog/ice-immigrationos-palantir-ai-track-immigrants/.
- [9] Kay, Jackie, et al. "Epistemic Injustice in Generative AI." Epistemic Injustice in Generative AI, Cornell University, 21 Aug. 2024, arxiv.org/html/2408.11441v1.
- [10] Koebler, Jason. "Inside a Powerful Database ICE Uses to Identify and Deport People." 404 Media, 9 Apr. 2025, www.404media.co/inside-a-powerful-database-ice-uses-to-identify-and-deport-people/.
- [11] Leah Hope Ajmani, Jasmine C. Foriest, Jordan Taylor, Kyle Pittman, Sarah Gilbert, and Michael Ann DeVito. 2024. Whose Knowledge is Valued? Epistemic Injustice in CSCW Applications. Proc. ACM Hum.-Comput. Interact. 8, CSCW2, Article 523 (November 2024), 28 pages. <https://doi.org/10.1145/3687062>
- [12] Lepri, B., Staiano, J., Sangokoya, D., Letouzé, E., Oliver, N. (2017). The Tyranny of Data? The Bright and Dark Sides of Data-Driven Decision-Making for Social Good. In: Cerquitelli, T.,

Quercia, D., Pasquale, F. (eds) Transparent Data Mining for Big and Small Data. Studies in Big Data, vol 32. Springer, Cham. https://doi.org/10.1007/978-3-319-54024-5_1

- [13] Rinaldi A, Teo SA. The Use of Artificial Intelligence Technologies in Border and Migration Control and the Subtle Erosion of Human Rights. *International and Comparative Law Quarterly*. 2025;74(1):61-89. doi:10.1017/S0020589325000090
- [14] Lorde, A. (2007). The master's tools will never dismantle the master's house. In *Sister outsider: Essays and speeches* (pp. 110-114). Crossing Press. (Original work published 1979)
- [15] Michael Muller and Angelika Strohmayer. 2022. Forgetting Practices in the Data Sciences. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (CHI '22)*. Association for Computing Machinery, New York, NY, USA, Article 323, 1–19. <https://doi.org/10.1145/3491102.3517644>
- [16] Sunstein, C.R. (2019). Algorithms, Correcting Biases. *Social Research: An International Quarterly* 86(2), 499-511. <https://dx.doi.org/10.1353/sor.2019.0024>.